




Intro to GitLab CI/CD

For teams Getting Started with GitLab CI/CD



Agenda



- What is CI/CD?
- GitLab CI/CD Overview
- GitLab CI/CD Setup
- GitLab CI/CD Runners
- Q&A



What is GitLab?



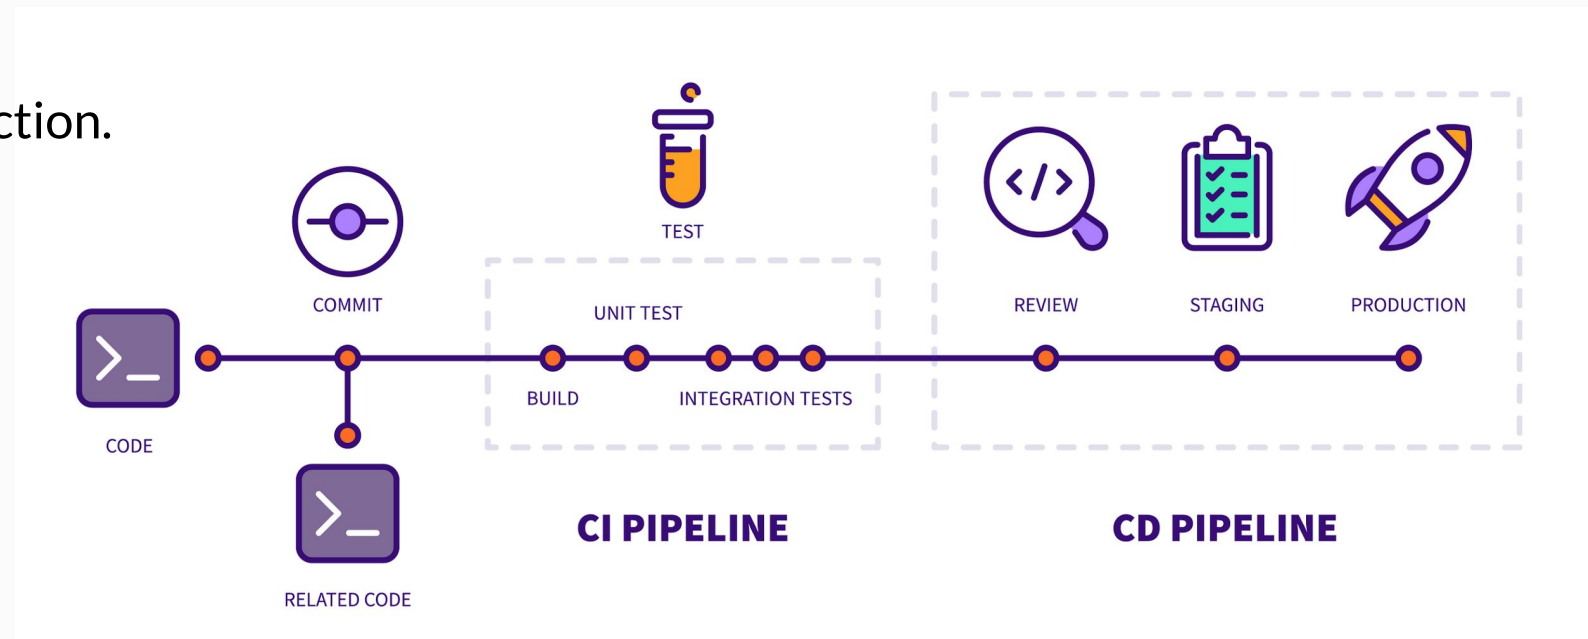
- DevOps Lifecycle Tool
 - Bridge dev and Ops.
- Git Repository Manager
 - Integrate code provided by your team in a shared repository.
- CI/CD
 - Empowers all teams to work together efficiently.
 - Powerful, scalable, end-to-end automation.



What is CI/CD?



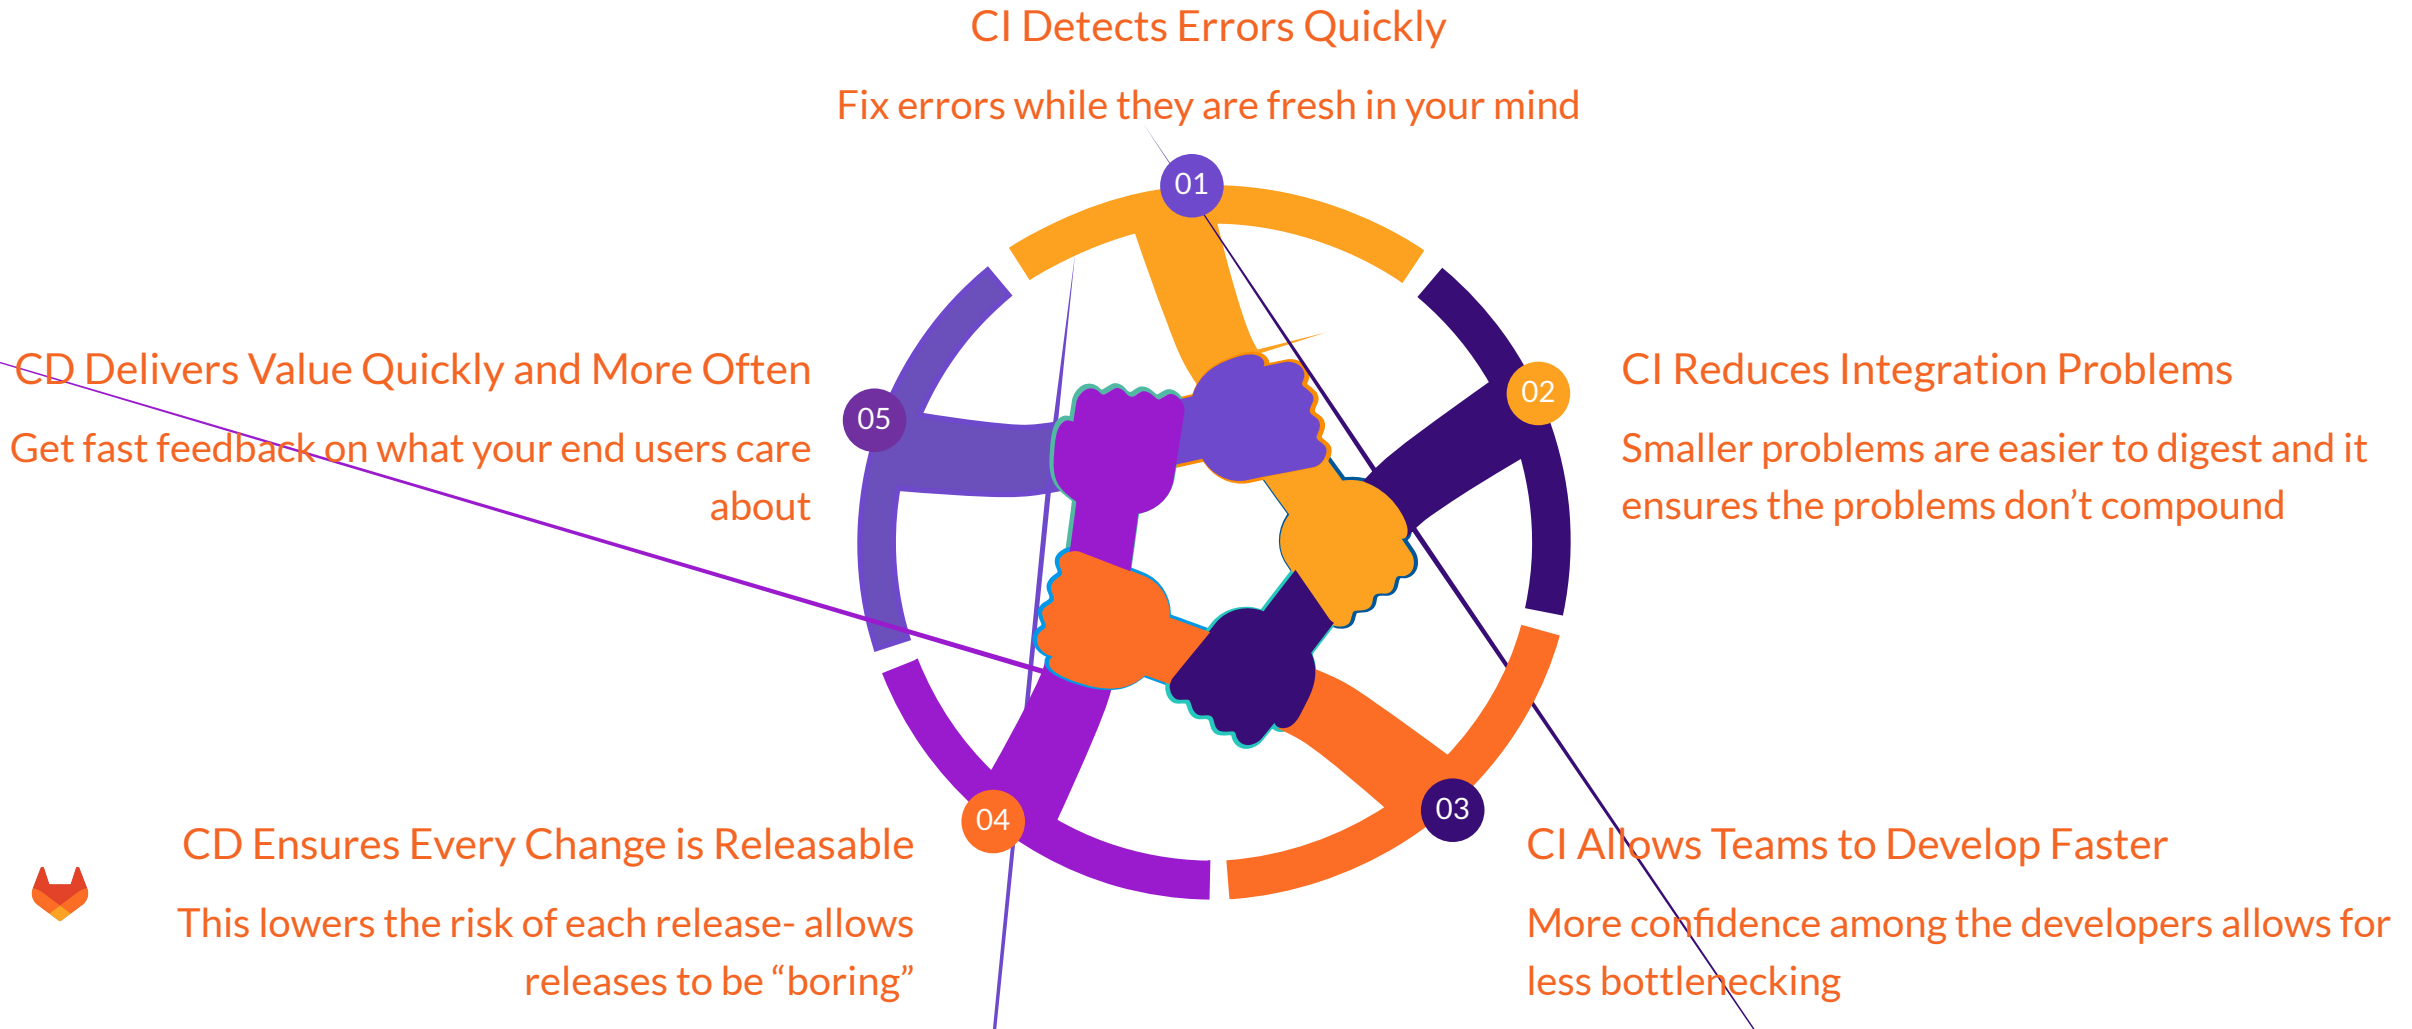
- Continuous Integration (CI)
 - Integrate code provided by your team in a shared repository.
- Continuous Delivery
 - Software released to production automatically.
- Continuous Deployment
 - Pushes changes to production.



Why use CI/CD?



CI/CD encourages collaboration across all departments and makes code creation and management easy, as well as provides the following specific benefits.



GitLab Recommended Process



Manage



Plan



Create



Verify



Package



Secure



Release



Configure



Monitor



Protect



Epics



Milestones



Issues



Assign Issue

Create Merge Request

Push Code

Automated Build / Test

Scan

Collaboration & Review

Approval

Review App

Push Fixes

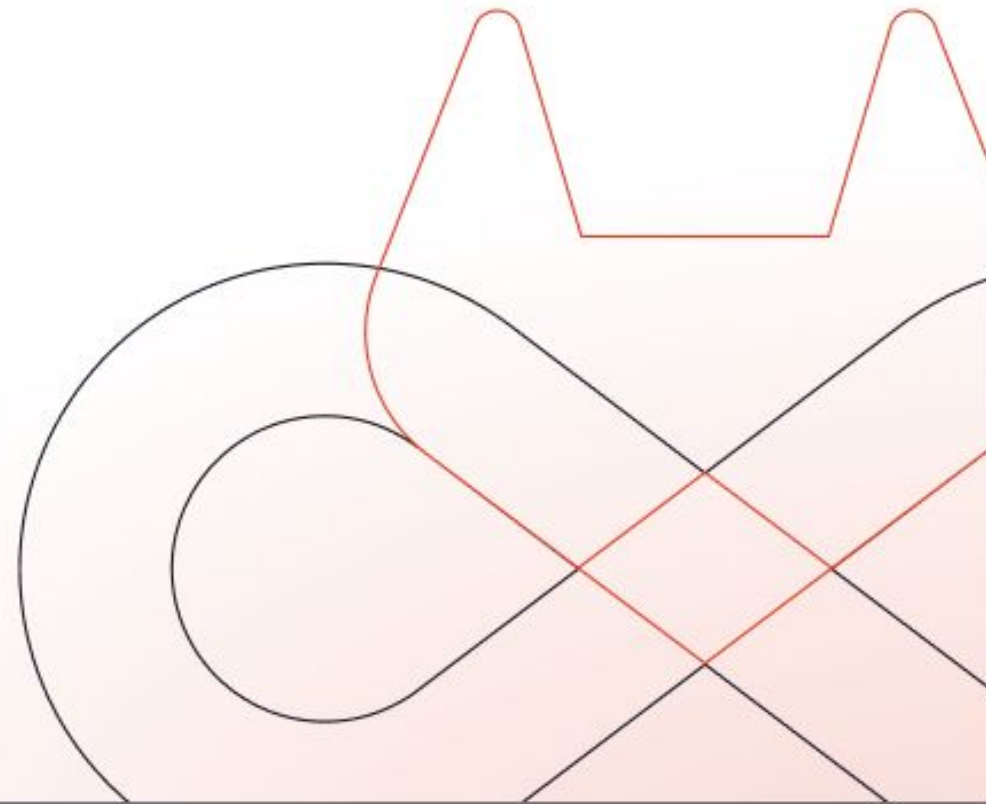
Merge Accepted

Release

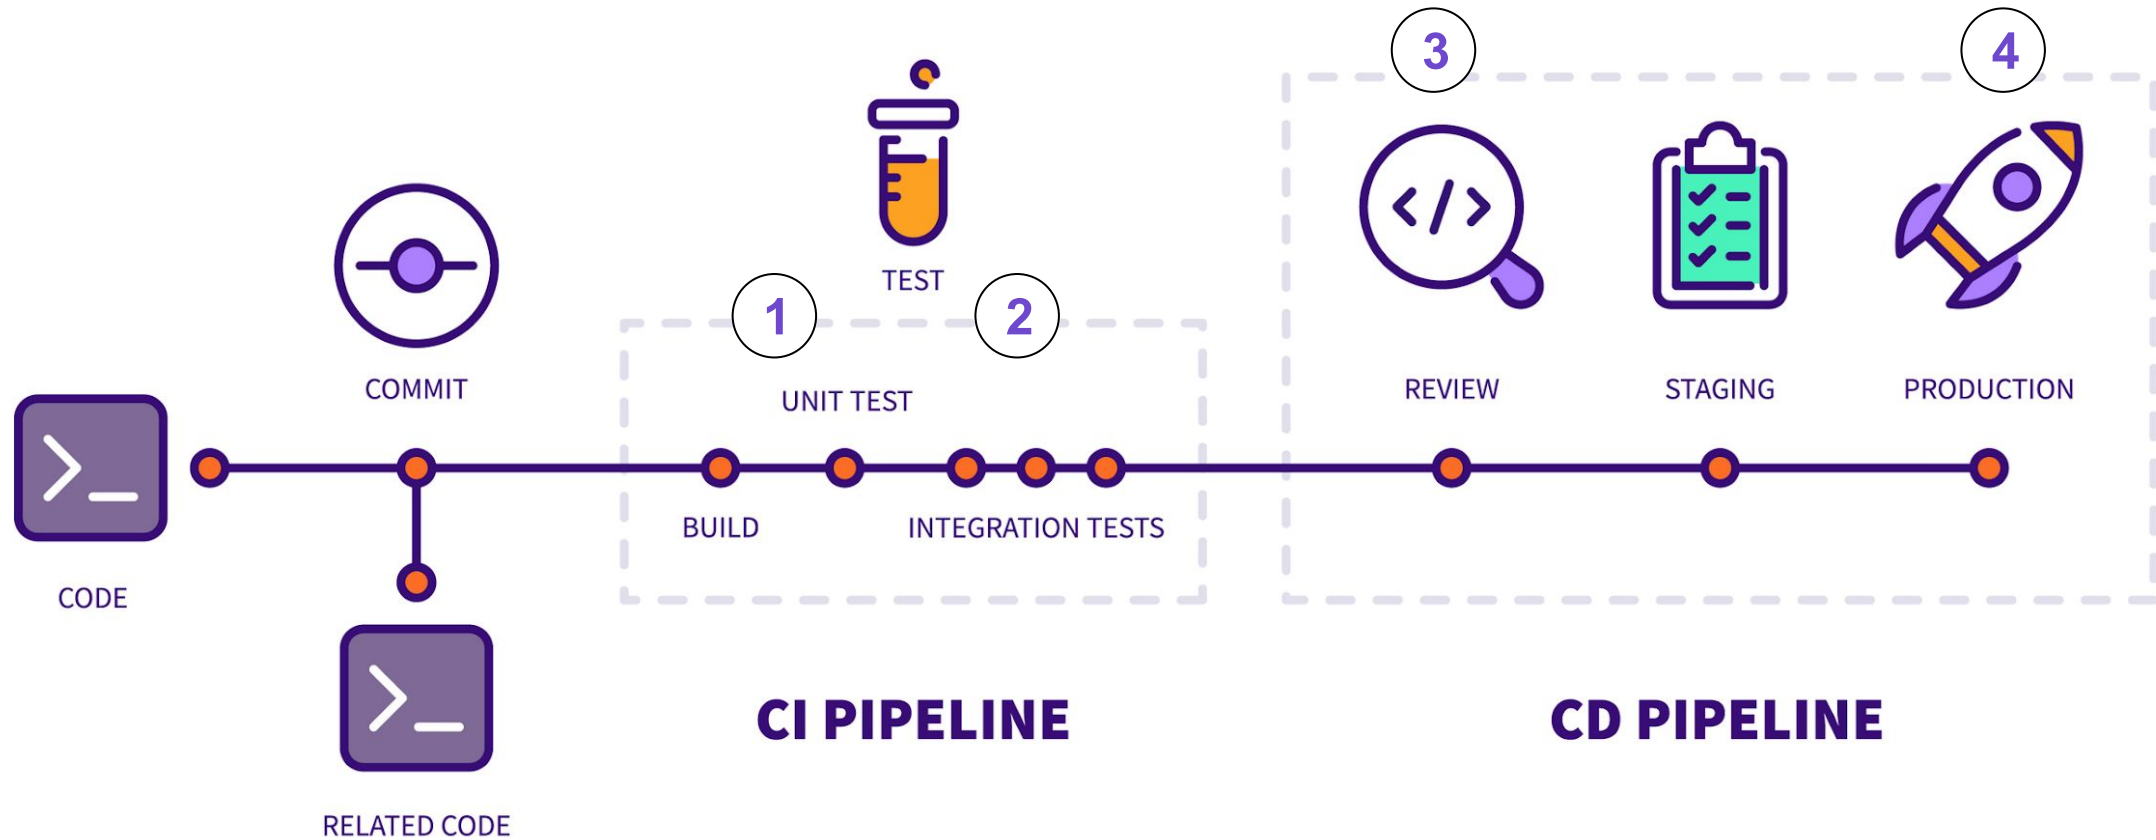
Deploy



GitLab CI/CD Overview



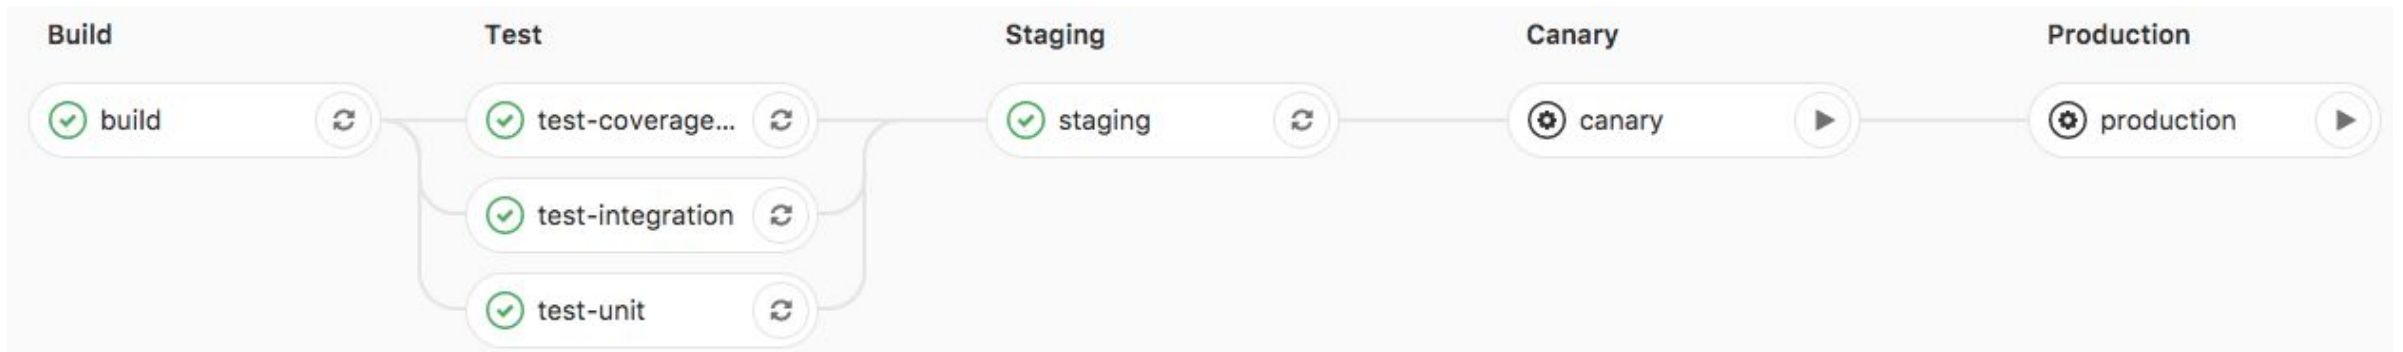
Anatomy of a CI/CD Pipeline



GitLab Pipeline Graph



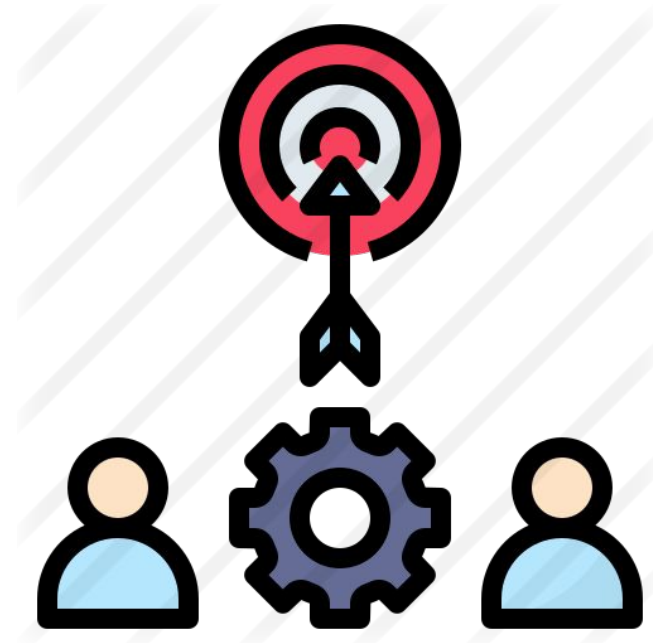
- Jobs define what we want to accomplish in our pipeline.
 - Executed by Runners
 - Executed in Stages
- Stages define when and how to run jobs.
 - Stages that run tests after stages that compile the code.
- Jobs in each stage are executed in parallel
 - If *all* jobs in a stage succeed, the pipeline moves on to the next stage.
 - if one job in a stage fails, the next stage is not (usually) executed




Ways to trigger GitLab pipeline



- Push your code to GitLab repository*
- Run it manually from the UI
- Schedule it to run at later time
- “Trigger”ed by upstream pipeline
- Use API to launch a pipeline with “trigger”



GitLab CI/CD Set-up



.gitlab-ci.yml Example



```
image:
  registry.gitlab.com/gitlab-examples/kubernete
  s-deploy
```

```
stages:
  - build
  - deploy
```

```
variables:
  KUBE_DOMAIN: example.com
```

```
build:
  stage: build
  script:
    - command build
  only:
    - main
```

```
deploy:
  stage: deploy
  script:
    - command deploy
  environment:
    name: production
    url: http://production.example.com
  variables:
    DISABLE_POSTGRES: "yes"
  only:
    - main
```



Build



build

Deploy



deploy



- A job is defined as a list of keywords that define the job's behavior.
- Configuration options for your GitLab .gitlab-ci.yml file.
- The keywords available for jobs are:
 - <https://docs.gitlab.com/ee/ci/yaml/>



- `image`
- `services`
- `script`
- `before_script` & `after_script`
- `variables`
- `Environment`
- `cache`
- `artifacts`
- `rules`
- `tags`
- `when`

Stages

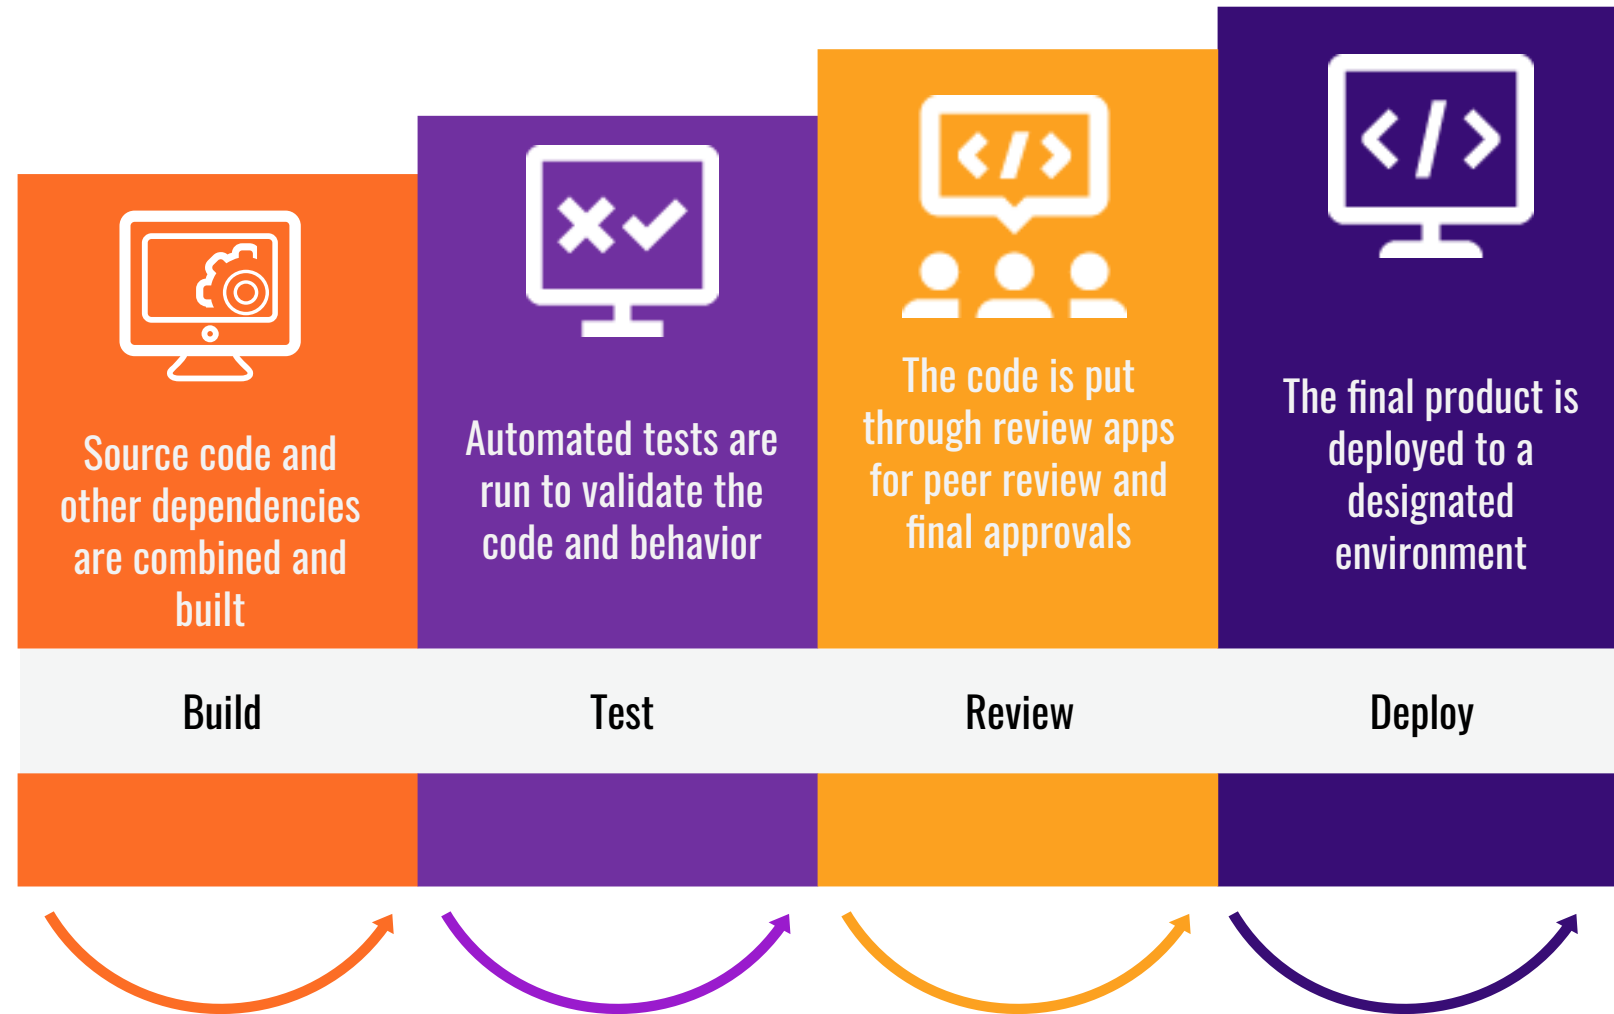


Default Stages: Build, Test, Review, & Deploy

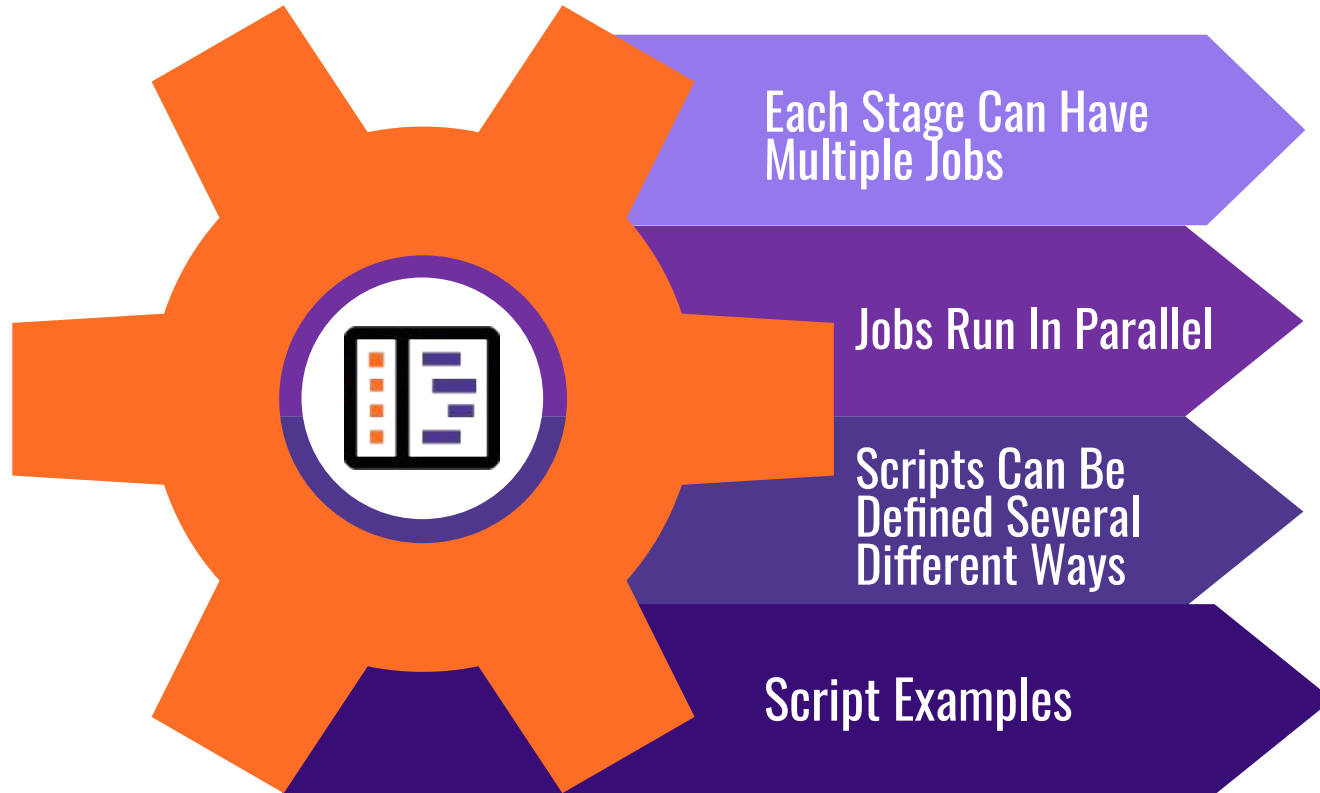
User can define custom stages & any number of jobs per stage

stages:

- build
- test
- review
- deploy



Stages separate jobs into logical sections while **Jobs** perform the actual tasks



```
build-code:  
  stage: build  
  script: build-it.sh
```

```
build-other-code:  
  stage: build  
  script: src/other/code/build-it.sh
```

```
script: command build
```

```
script:  
  - npm install  
  - npm build
```

```
script:  
  - scripts/build_script.sh
```



Basic Parameters



```
test:
  script:
    - apt-get update -qy
    - bundle install --path /cache
    - bundle exec rake test

staging:
  stage: deploy
  script:
    - gem install dpl
    - dpl --provider=heroku --app=ruby-test-staging --api-key=$HEROKU_KEY
  only:
    - main

production:
  stage: deploy
  script:
    - gem install dpl
    - dpl --provider=heroku --app=ruby-prod --api-key=$HEROKU_PROD_KEY
  only:
    - tags
```



Image



Images are pulled from Docker Hub by default

Use of a public image:

```
image: ruby:2.3
```

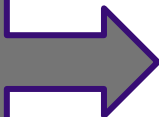
Images stored in the GitLab Container Registry

Use of a custom image:

```
image:  
'registry.gitlab.com/gitlab-org/ci-training-sample:latest'
```



.gitlab-ci.yml
build so far



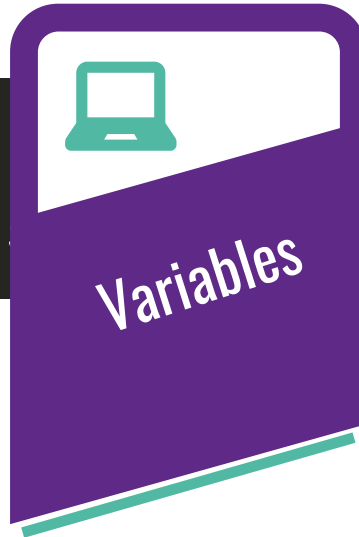
```
image: registry.example.com/k8-deploy:latest
```

Services & Variables



```
services:  
- postgres
```

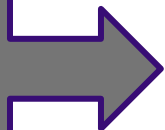
Services lines tell the Runner that additional images are needed



```
variables:  
- POSTGRES_DB:  
rails-sample-1_test  
- POSTGRES_USER: root  
POSTGRES_PASSWORD:
```

Variables also defined in Project > Settings > CI/CD > Variables

.gitlab-ci.yml
build so far



```
image: registry.example.com/k8-deploy:latest  
services:  
- postgres  
variables:  
- POSTGRES_DB: rails-sample-1_test
```

What Our .gitlab-ci.yml looks like so far...



```
image: registry.example.com/k8-deploy:latest
services:
  - postgres
variables:
  - POSTGRES_DB: rails-sample-1_test
stages:
  - build
  - test
  - deploy
deploy-code:
  stage: deploy
  script:
    - command deploy
```



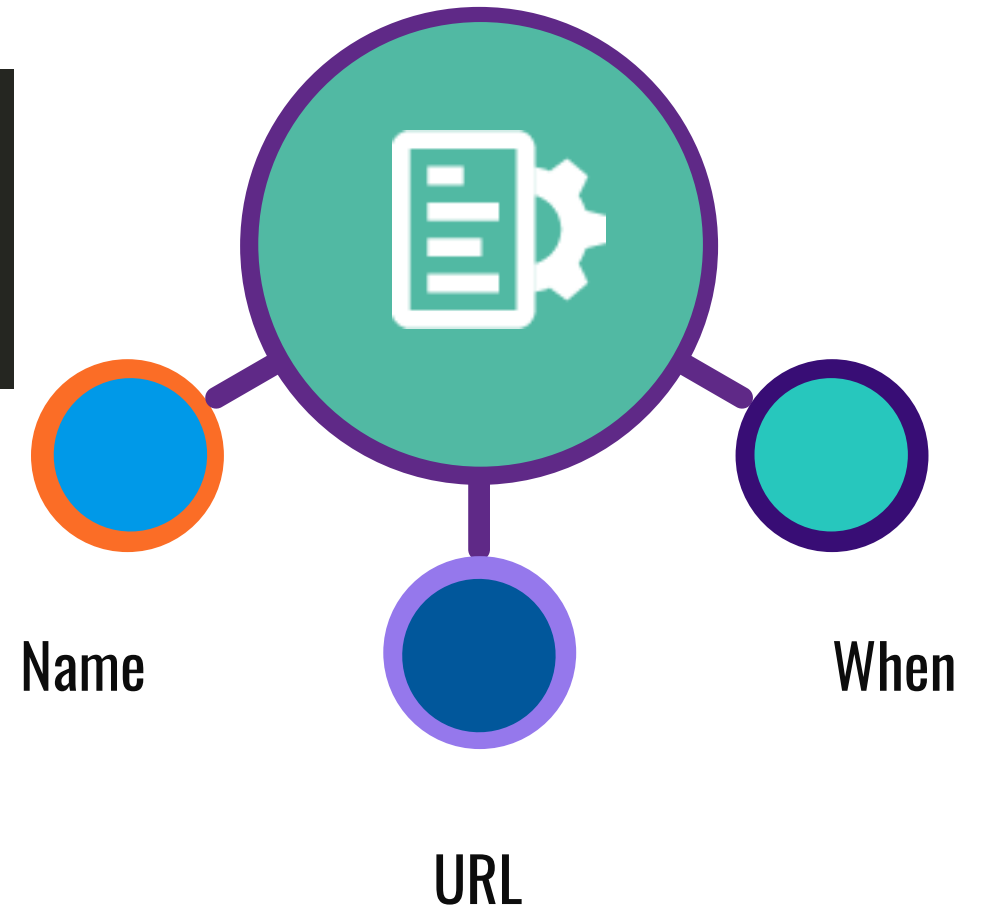
Environments



The environment keyword defines where the app is deployed and is defined by 3 parts.

```
environment:  
  name: prod  
  url: http://$CI_PROJECT_NAME.$KUBE_DOMAIN  
  when: manual
```

When triggers jobs & stages manually (e.g. deploy to production)



Only & Except- Restricting When a Job is Executed

```
pseudo-deploy:  
  stage: deploy  
  script:  
  - command deploy_review  
  only:  
  - branches  
  except:  
  - main  
  environment:  
    name: review  
    url: http://$CI_PROJECT_NAME-review.$KUBE_DOMAIN
```



Only

The name of branch
to execute on (in
this case all
branches)



Except

Branches NOT to
execute on with
exception to the
main Branch

The [rules syntax](#) is an improved, more powerful solution for defining when jobs should run or not. *Consider using rules instead of only/except to get the most out of your pipelines.*



Rules - Restricting When a Job is Executed



```
pseudo-deploy:  
  stage: deploy  
  script:  
    - command deploy_review  
  rules:  
    - if: '$CI_COMMIT_REF_NAME == "main"'  
      when: never  
    - when: always  
  environment:  
    name: review  
    url: http://$CI_PROJECT_NAME-review.$KUBE_DOMAIN
```





Run before and after the script defined in each job

- Can update the image with the latest version of components
- They run within the job and can interact with the job



before_script

is used to define a command that should be run before each job, including deploy jobs, but after the restoration of any artifacts

```
before_script:  
- echo $CI_BUILD_STAGE  
- apt-get update  
- apt-get install node-js -y  
- bundle install  
- npm install  
after_script:  
- rm temp/*.tmp
```

after_script

is used to define the command that will be run after each job, including failed ones.



Cache & Artifacts



Cache is used to pass information between jobs & stages by storing project dependencies

```
cache:  
  paths:  
    - binary/  
    - .config
```

There may be build artifacts you want to save

```
artifacts:  
  when: on_success  
  paths:  
    - bin/target
```



What Our .gitlab-ci.yml looks like so far...



```
image: registry.example.com/k8-deploy:latest
services:
  - postgres
variables:
  - POSTGRES_DB: rails-sample-1_test
cache:
  paths:
    - binary/
stages:
  - build
  - test
  - deploy
deploy-code:
  stage: deploy
  script:
    - command deploy
  environment:
    name: production
    url: http://$CI_PROJECT_NAME.$KUBE_DOMAIN
  when: manual
  only:
    - main
```



```
build-it:
  stage: build
  script:
    - command build
  only:
    - main
  artifacts:
    when: on_success
    paths:
      - bin/target
```



- Tags are used to select a specific runner
 - CI tags are different from Git tags
- Runners with the required tags can pick-up the job
 - If a Runner has more tags than required, it can still run that particular job; including if the job requires no tags at all

```
job-name:  
  tags:  
    - ruby  
    - test
```



What Our .gitlab-ci.yml looks like so far...

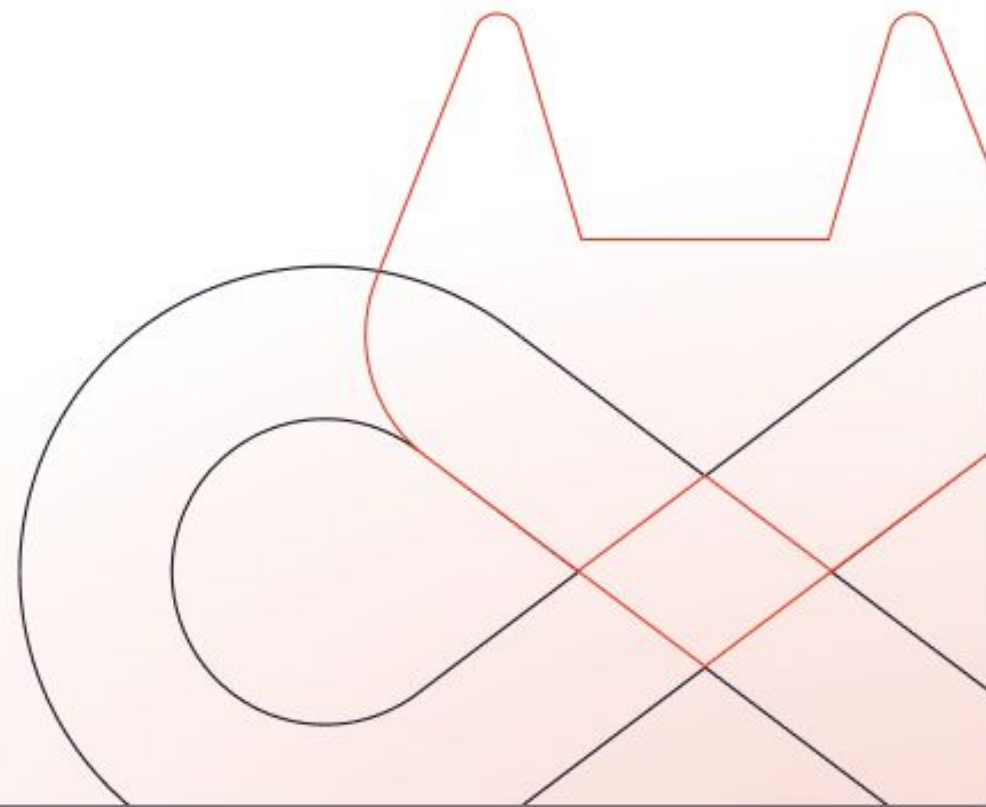


```
image: registry.example.com/k8-deploy:latest
services:
  - postgres
variables:
  - POSTGRES_DB: rails-sample-1_test
cache:
  paths:
    - binary/
stages:
  - build
  - test
  - deploy
deploy-code:
  stage: deploy
  script:
    - command deploy
  environment:
    name: production
    url: http://$CI_PROJECT_NAME.$KUBE_DOMAIN
  when: manual
  only:
    - main
```



```
build-it:
  stage: build
  script:
    - command build
  only:
    - main
  tags:
    - osx
    - ios
  artifacts:
    when: on_success
    paths:
      - bin/target
```

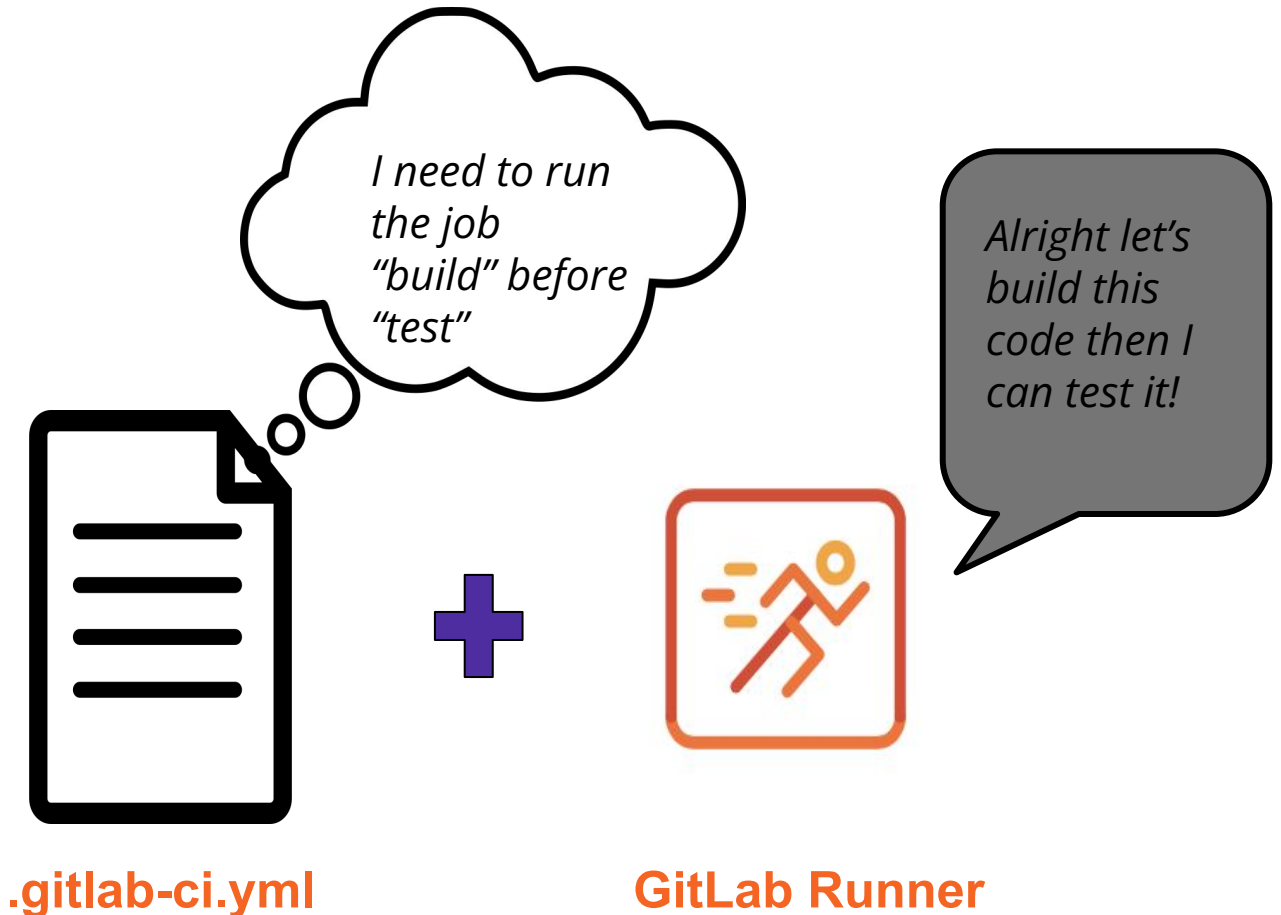
GitLab CI/CD Runners



Configuration File + Runner

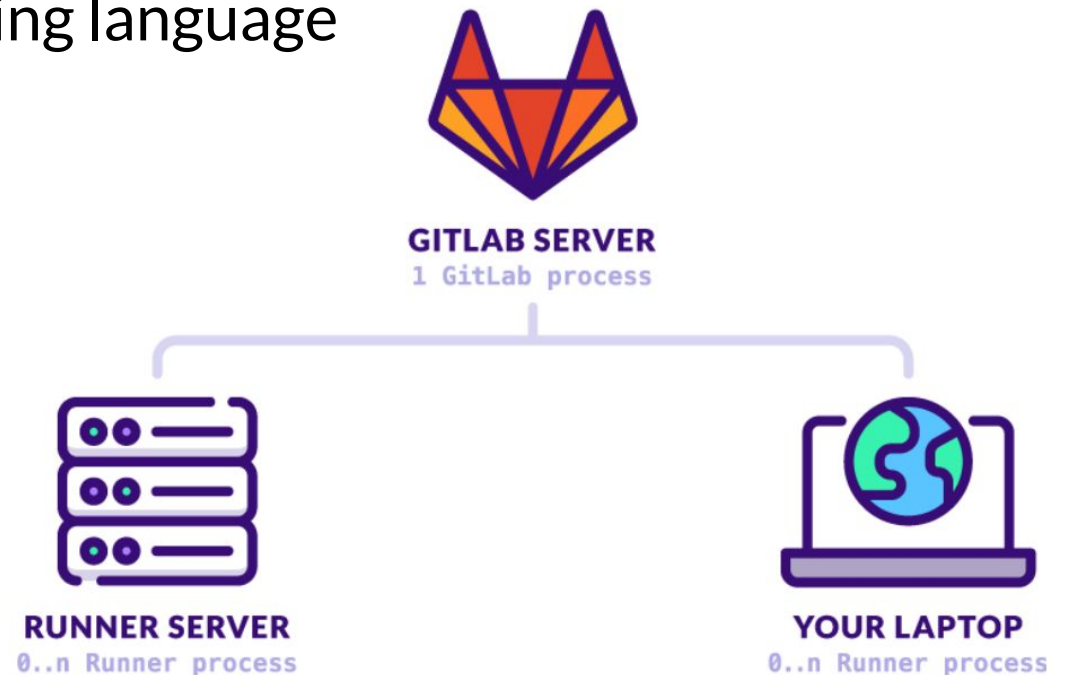


- .gitlab-ci.yml file
 - Instructions for GitLab CI/CD jobs.
 - Lives in the root of the repository
- GitLab Runner
 - Lightweight agent that runs CI/CD jobs.





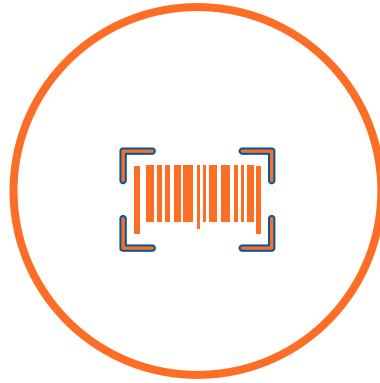
- The GitLab runner can be installed on any platform where you build Go binaries.
 - Linux, macOS, Windows, FreeBSD, Cloud Provider, Bare Metal, Your work station and Docker
- The GitLab runner can test any programming language
 - .Net, Java, Python, C, PHP and others.
- Created by an Administrator



A Runner Can Be....



Shared or Specific



Tagged or Untagged



Protected or Not Protected



Shared vs. Specific Runners



Shared Runners

Available to every project with similar requirements

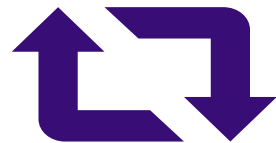
Description



Included in the pool for all projects



Managed by GitLab Admin



Typically auto-scaling or otherwise scaled

VS



Specific Runners

Tied to one or more specific projects

Description



In the pool for ONLY specific projects



Managed by Runner Owner(s)



Typically for specialized builds, or if an org needs to do so for billing



Tagged vs. Untagged



● babb8003

Pause

Disable for this project

WIN-2012-EC2

#244878

windows

```
msbuild:
  stage: 📦 build
  script:
    - cd csharp-msbuild
    - buildit
  artifacts:
    paths:
      - HelloWorld.exe
  tags:
    - windows
```

● edb9fc6c

Pause

Disable for this project

Brendans-MacBook-Pro.local

#210789

```
java:spring-boot:
  stage: 📦 build
  image: maven:3.5-jdk-8-slim
  script:
    - cd java
    - cd spring-boot
    - mvn package
  artifacts:
    paths:
      - java/spring-boot/target/*
```

Tagged

Only used to run jobs tagged with same tag

Untagged

Used to run jobs with no tags

Protected vs. Non-Protected



Protected

Characteristics

- ONLY runs jobs from
- Protected Branches
 - Protected Tags

Typically used for runners containing deploy keys or other sensitive capabilities



Non-Protected

Characteristics

- Runs jobs from ANY branch
- Used for ANY build

Additional Runner Options



Runner #1323

- Active

☒ Paused Runners don't accept new jobs
- Protected

☐ This runner will only run on pipelines triggered on protected branches
- Run untagged jobs

☒ Indicates whether this runner can pick jobs without tags
- Lock to current projects

☒ When a runner is locked, it cannot be assigned to other projects

IP Address 72.195.135.57

Description MacBook-Pro.local

Maximum job timeout

This timeout will take precedence when lower than project-defined timeout and accepts a human readable time input language like "1 hour". Values without specification represent seconds.

Tags

You can set up jobs to only use Runners with specific tags. Separate tags with commas.

Save changes

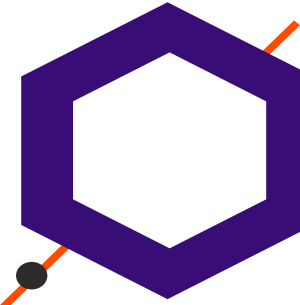


Executors: Common




Shell

Directly run commands as if writing them into terminal
(bash or sh) or command prompt (cmd) or powershell




Docker Machine

“Main” machine scales up runners with
any executor on demand
Typical in cloud deployments



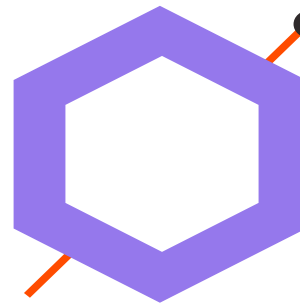
Docker

Execute inside of a docker image
Most common!



Kubernetes

Runs as a pod in a K8s cluster
Can also feature auto-scaling





1. VirtualBox

Base VM for runner
“Main” creates a new VM for each needed runner



2. Parallels

Hint: Parallels is a nice platform on top of
VirtualBox



3. SSH

Similar to shell, but not as many features (bash only, no caching)
Does allow you to SSH and execute commands on a machine you
might not want to install runner on





Q&A

